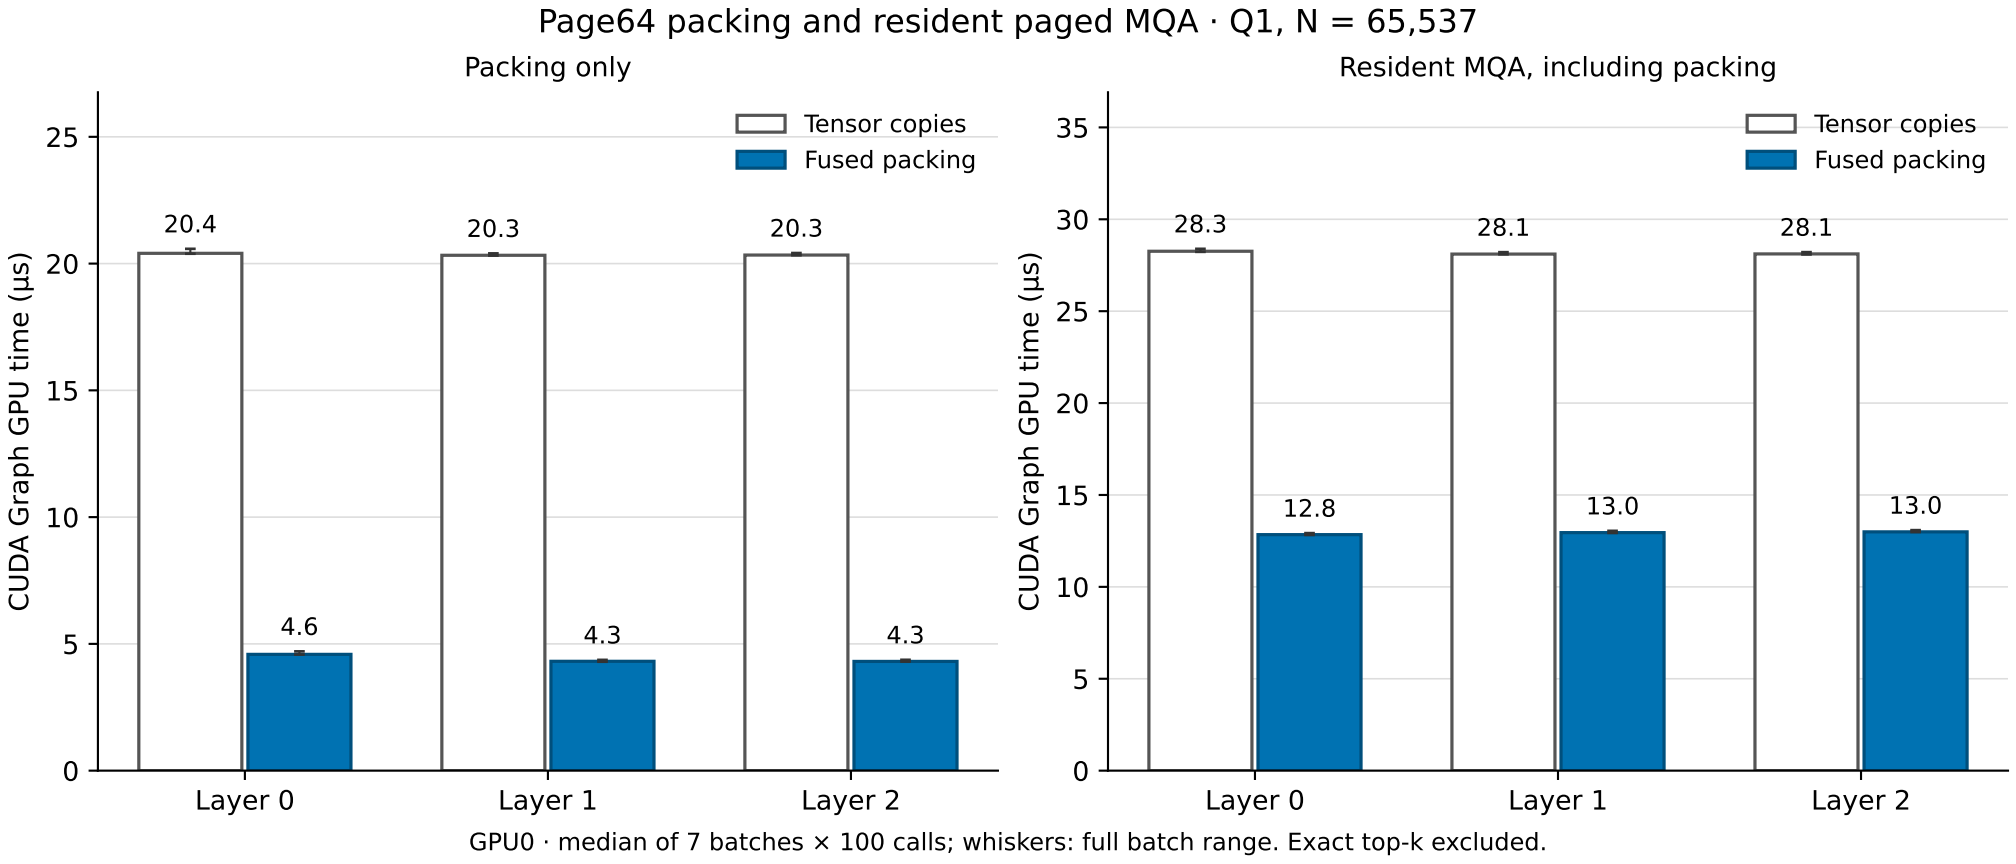
Page64 packing and resident paged MQA · Q1, N = 65,537
Packing only
Resident MQA, including packing
Tensor copies
Fused packing
CUDA Graph GPU time (µs)
0
5
10
15
20
25
30
35
20.4
20.3
20.3
4.6
4.3
4.3
28.3
28.1
28.1
12.8
13.0
13.0
Layer 0
Layer 1
Layer 2
GPU0 · median of 7 batches × 100 calls; whiskers: full batch range. Exact top-k excluded.