


DeepSeek V3.2 | History prefill: one-layer timeline

Measured Nsight schedule · H200 / SM90 · Q=1,024 · chunk 63 · tokens 64,512-65,535

Prefix before chunk: 64,512 tokens · final history H=65,536 · pool P=65,536

Physical block L1 (zero-based) · cold request 0 · each row starts at this chunk's L0 GPU finish · same time scale



No history KV H2D in these windows; ECHO uses its resident indexer. Small control H2D copies remain in Aux / cache.

All intersecting GPU activities and gaps are retained, including adjacent-layer asynchronous work. Writeback labels identify its owning layer.

One instrumented cold-request sample; these windows are not formal layer latencies. deepseek_mfu_c10_profile_20261006_02