


DeepSeek V3.2 | One-layer computation and communication

Measured Nsight schedule · H200 / SM90 · candidate Q=128 · history H=65,536 · pool P=65,536

Block L1 (zero-based) · revisit request 16 · each row starts at L0 GPU finish · same time scale



Logical lanes preserve actual GPU timestamps. KV gathers read mapped host memory on SMs; they are not copy-engine DMA.

White gaps are retained. Fused-kernel internals are not resolved. One instrumented sample; these windows are not benchmark layer latencies.